

Deep Learning in Data Science

Harald Stiff Erik von Keyserlingk Jacob Stuart

Image Generation Using Deep Convolutional Generative Adversarial Networks

Problem Formulation

Background

Technical implementation

- Generator

- Discriminator

Results

- Generated images

- Latent space exploration

Discussion

Problem Formulation

Background

Technical implementation

Generator

Discriminator

Results

Generated images

Latent space exploration

Discussion

- Generate images resembles the training set
- Generate images that are diverse and covers the features of the whole training set
- Investigate the mapping from latent space to generated image space

Our work and models are based on (Radford et al., 2015)
*Unsupervised Representation Learning With Deep Convolutional
Generative Adversarial Networks*

- Simple guidelines for good results
- Stability during training
- Accepted as state of art for GAN architecture designs

Tools used:

- Python3
- Tensorflow
- Keras
- Numpy
- Google computation was used for running the code

Problem Formulation

Background

Technical implementation

Generator

Discriminator

Results

Generated images

Latent space exploration

Discussion

Composed of two models:

$$\underbrace{G(z, \Theta_g)}_{\text{Generator}}$$

&

$$\underbrace{D(x, \Theta_d)}_{\text{Discriminator}}$$

where

$$\underbrace{\Theta_g = \{\theta_g^i\}_{i=1}^n}$$

Generator parameters

&

$$\underbrace{\Theta_d = \{\theta_d^j\}_{j=1}^m}$$

Discriminator parameters

Composed of two models:

$$\underbrace{G(z, \Theta_g)}_{\text{Generator}}$$

&

$$\underbrace{D(x, \Theta_d)}_{\text{Discriminator}}$$


where

$$\underbrace{\Theta_g = \{\theta_g^i\}_{i=1}^n}_{\text{Generator parameters}}$$


&

$$\underbrace{\Theta_d = \{\theta_d^j\}_{j=1}^m}_{\text{Discriminator parameters}}$$

Generator: Maps an input vector $z \sim p_z(z)$ to a tensor x_g .



Discriminator: Maps tensors $\mathbf{x}_{real} \sim p_{data}(\mathbf{x})$ and \mathbf{x}_g to a scalar score c .



Discriminator update:

$$\underbrace{\mathcal{B}_d}_{\text{Batch}} = \left\{ (\mathbf{x}_{real}^i, y = 1), (\mathbf{x}_g^i, y = 0) \right\}_{i=1}^{N/2}$$

$$\underbrace{\mathcal{L}(\mathcal{B}_d, \Theta_d)}_{\text{Loss}} = - \frac{1}{|\mathcal{B}_d|} \sum_{(\mathbf{x}, y) \in \mathcal{B}_d} y \log(D(\mathbf{x})) + (1 - y) \log(1 - D(\mathbf{x}))$$

$$\underbrace{\frac{\partial \mathcal{L}(\mathcal{B}_d, \Theta_d)}{\partial \theta}}_{\text{Gradients}}, \quad \forall \theta \in \Theta_d$$

Discriminator update:

$$\underbrace{\mathcal{B}_d}_{\text{Batch}} = \left\{ (\mathbf{x}_{real}^i, y = 1), (\mathbf{x}_g^i, y = 0) \right\}_{i=1}^{N/2}$$

$$\underbrace{\mathcal{L}(\mathcal{B}_d, \Theta_d)}_{\text{Loss}} = - \frac{1}{|\mathcal{B}_d|} \sum_{(\mathbf{x}, y) \in \mathcal{B}_d} y \log(D(\mathbf{x})) + (1 - y) \log(1 - D(\mathbf{x}))$$

$$\underbrace{\frac{\partial \mathcal{L}(\mathcal{B}_d, \Theta_d)}{\partial \theta}}_{\text{Gradients}}, \quad \forall \theta \in \Theta_d$$

Discriminator update:

$$\underbrace{\mathcal{B}_d}_{\text{Batch}} = \left\{ (\mathbf{x}_{real}^i, y = 1), (\mathbf{x}_g^i, y = 0) \right\}_{i=1}^{N/2}$$

$$\underbrace{\mathcal{L}(\mathcal{B}_d, \Theta_d)}_{\text{Loss}} = - \frac{1}{|\mathcal{B}_d|} \sum_{(\mathbf{x}, y) \in \mathcal{B}_d} y \log(D(\mathbf{x})) + (1 - y) \log(1 - D(\mathbf{x}))$$

$$\underbrace{\frac{\partial \mathcal{L}(\mathcal{B}_d, \Theta_d)}{\partial \theta}}_{\text{Gradients}}, \quad \forall \theta \in \Theta_d$$

Generator update:

$$\underbrace{\mathcal{B}_g}_{\text{Batch}} = \left\{ (z_i, y = 1) \right\}_{i=1}^N$$

$$\underbrace{\mathcal{L}(\mathcal{B}_g, \Theta_d, \Theta_g)}_{\text{Loss}} = -\frac{1}{|\mathcal{B}_g|} \sum_{(z,y) \in \mathcal{B}_g} y \log D(G(z))$$

$$\underbrace{\frac{\partial \mathcal{L}(\mathcal{B}_g, \Theta_d, \Theta_g)}{\partial \theta}}_{\text{Gradients}}, \quad \forall \theta \in \Theta_g$$

Generator update:

$$\underbrace{\mathcal{B}_g}_{\text{Batch}} = \left\{ (z_i, y = 1) \right\}_{i=1}^N$$

$$\underbrace{\mathcal{L}(\mathcal{B}_g, \Theta_d, \Theta_g)}_{\text{Loss}} = -\frac{1}{|\mathcal{B}_g|} \sum_{(z,y) \in \mathcal{B}_g} y \log D(G(z))$$

$$\underbrace{\frac{\partial \mathcal{L}(\mathcal{B}_g, \Theta_d, \Theta_g)}{\partial \theta}}_{\text{Gradients}}, \quad \forall \theta \in \Theta_g$$

Generator update:

$$\underbrace{\mathcal{B}_g}_{\text{Batch}} = \left\{ (z_i, y = 1) \right\}_{i=1}^N$$

$$\underbrace{\mathcal{L}(\mathcal{B}_g, \Theta_d, \Theta_g)}_{\text{Loss}} = -\frac{1}{|\mathcal{B}_g|} \sum_{(z,y) \in \mathcal{B}_g} y \log D(G(z))$$

$$\underbrace{\frac{\partial \mathcal{L}(\mathcal{B}_g, \Theta_d, \Theta_g)}{\partial \theta}}_{\text{Gradients}}, \quad \forall \theta \in \Theta_g$$

Problem Formulation

Background

Technical implementation

- Generator

- Discriminator


Results

- Generated images

- Latent space exploration

Discussion


Generator



Generator architecture:

- Input: Noise vector 100×1 , uniformly distributed white noise
- Deconvolutional layers
- Kernel size 5×5
- Fully connected first layer
- Activation function: LeakyReLU
- Activation function last layer: tanh

Discriminator



Discriminator architecture:

- Input: an image
- Convolutional layers
- Kernel size 5×5
- Activation function: LeakyReLU
- Fully connected last layer
- Activation function last layer: Sigmoid

Problem Formulation

Background

Technical implementation

- Generator

- Discriminator

Results

- Generated images

- Latent space exploration

Discussion



A FEW EPOCHS LATER

Results (MNIST)




Figure 2: MNIST dataset











Figure 3: Generated numbers



Results (CIFAR 10) 100 Epochs



Results (Flowers) 100 Epochs



Cats 100 Epochs



Results (Cats)




Figure 4: Generated images after 55 epochs of training.




Figure 5: Generated images after 60 epochs of training.

Latent space exploration (MNIST)



Figure 6: Latent space exploration for MNIST

Latent space exploration (Flowers & CIFAR 10)




Figure 7: Latent space exploration for flowers




Figure 8: Latent space exploration for CIFAR10

Problem Formulation

Background

Technical implementation

- Generator

- Discriminator

Results

- Generated images

- Latent space exploration

Discussion

Model collapse

When one of the networks outperforms the other

Our case the discriminator outperformed

Generator found a local minimum




Figure 9: Generated images after 55 epochs of training.




Figure 10: Generated images after 60 epochs of training.

Background variation

The background in cat images are diverse and does not have anything with the cats attribute to do.




Figure 11: Cat



Figure 12: Cat

DCGAN works well with MNIST (handwritten numbers)

The network works okay with images of flowers

Hard for the network to generate images with diverse backgrounds,
such as cats

Thank you for
your attention!




Figure 13: Summer flower